



WIZELINE

Cloud Native Serverless

#sre-emea



Your Questions during the session

Maybe in the session, we does not leave much time for questions.

Please post your questions to following Google doc

<https://drive.google.com/file/d/15PvTVs3EBNVrUvA9TOfy37ImZ4-fQuCz/view?usp=sharing>

OR let's discuss in the team Slack channel **#sre-emea**

Repo → [cloudnative-serverless-workshops](#)



Agenda

- About the Cloud native Serverless workshop series.
 - Quick review about previous session.
 - New stuff in the repository
- Getting Context
- Deployment Strategies
 - Rolling Update
 - Canary Deployment
 - Blue Green Deployment
- Serverless Deployment using [Knative Traffic management](#)



Getting Context



Workshop Knative Traffic Management

Steps required.

```
> make help
hacking          install tooling required, Kind , Kubectl , make ..
create-cluster   Create Kind Cluster
delete-cluster   Delete Kind Cluster
knative-install  Install all the knative components in the cluster
knative-uninstall UnInstall all the knative components in the cluster
knative-show     Knative configuration and resources
istio-install    Install the istio component in the cluster
kubeless-install Install all the kubeless components in the cluster
kubeless-uninstall UnInstall all the kubeless components in the cluster
monitoring-install Install monitoring Operator Stacks (Prometheus, Grafana)
monitoring-uninstall Install monitoring Operator Stacks (Prometheus, Grafana)
knative-serving-workshop-build Build Knative Helloworld-go servicing By Default
knative-serving-workshop-serve Run Knative Helloworld-go servicing By Default .
knative-bluegreen-workshop Run Knative bluegreen
```

<https://github.com/erasmodominguezdc/cloudnative-serverless-workshops/tree/deployments>



Deployment Considerations

proprietary + confidential

WIZELINE

www.wizeline.com

As software Engineers, we must:

- Minimize impact to consumers/clients.
- Provide a rollback and recovery strategy.
- Be fast and robust in the Deployment process.
- Do not affect and impact the whole system.



Not Many years ago....





Kubernetes deployment

*A Kubernetes deployment is basically a **resource object** in Kubernetes that defines the **desired state** of your application.*

Kubernetes is basically an API with a Database that convert our yaml manifest into objects that represent a desired state



Kubernetes Deployment

```
apiVersion: apps/v1beta1 #Kubernetes API
kind: Deployment
metadata:
  name: hello-deploy
spec:
  replicas: 10
  selector:
    matchLabels:
      app: hello-world
  minReadySeconds: 10
  strategy:
    type: Recreate /RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 1
```



Kubernetes deployment

proprietary + confidential


WIZELINE

www.wizeline.com

*Ok!! Kubernetes is pretty cool , with some YAML hacks
I can define my deployments as I need it, but....I just
need to **manage the traffic!!***



Knative Serving Architecture



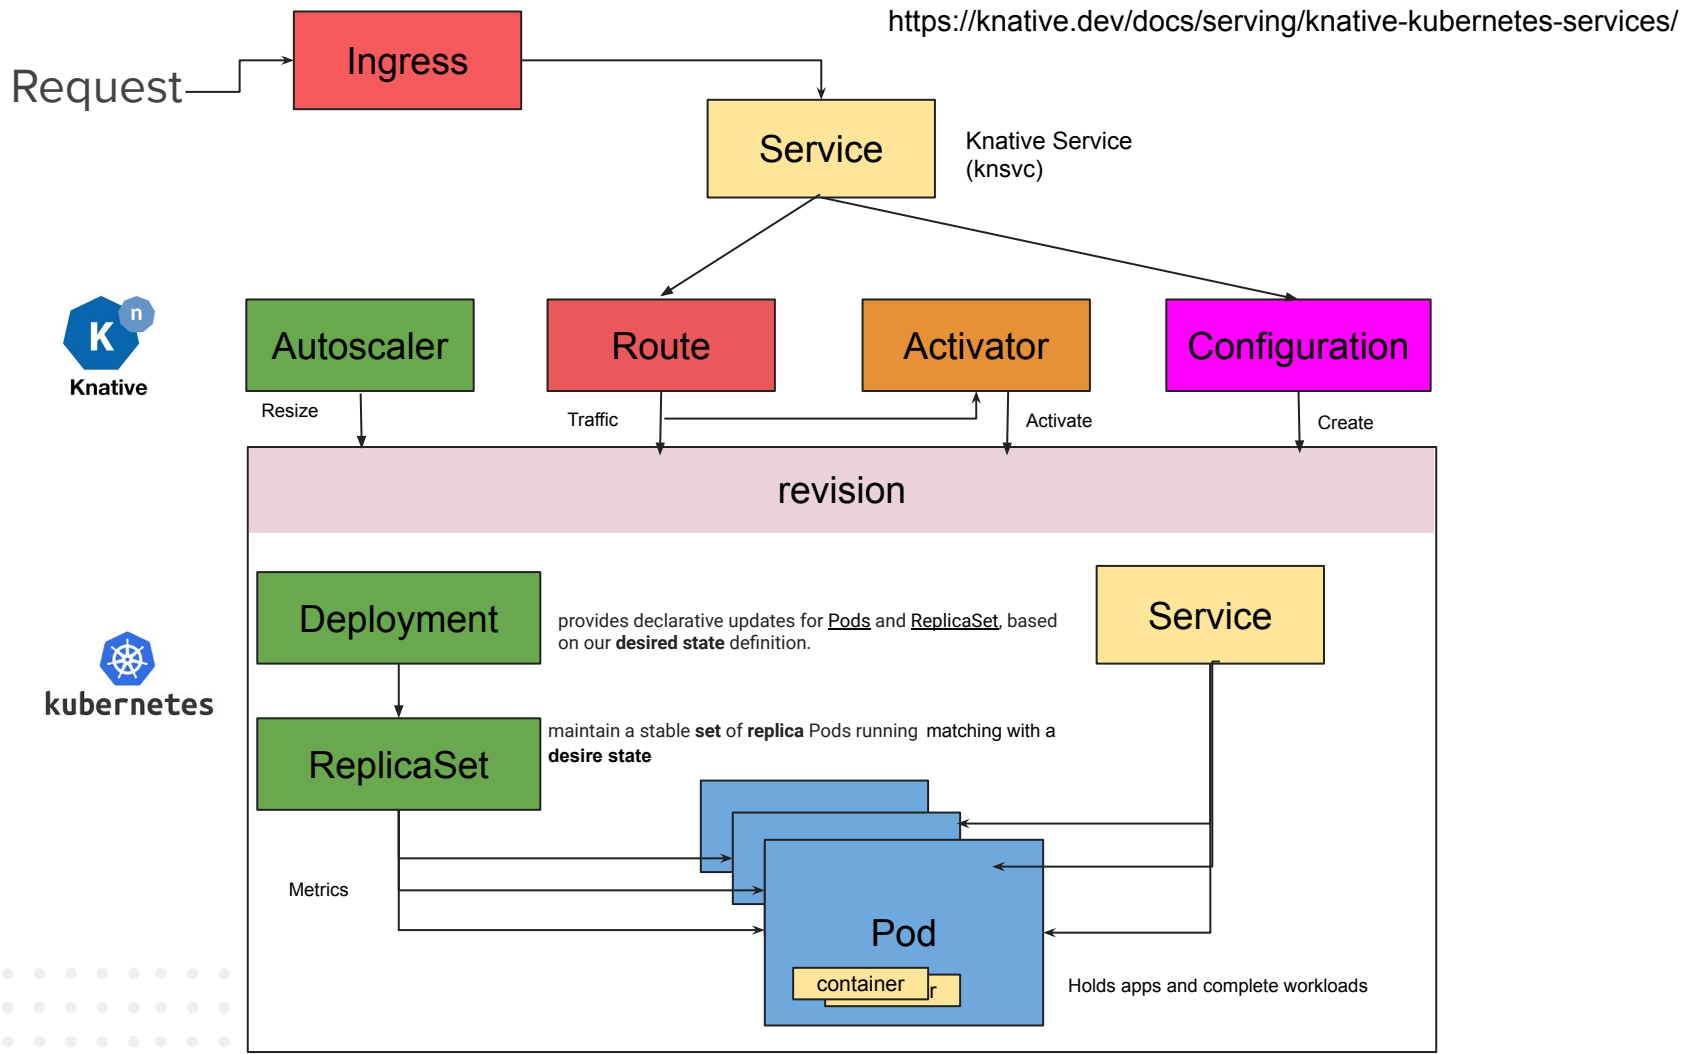
<https://knative.dev/docs/serving/>



Knative



kubernetes





Knative Traffic Management

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: foo
spec:
  template:
    # removed for brevity
  traffic: ❶
  - tag: v1 ❷
    revisionName: foo-v1 ❸
    percent: 50 ❹
  - tag: v2
    revisionName: foo-v2
    percent: 50
```

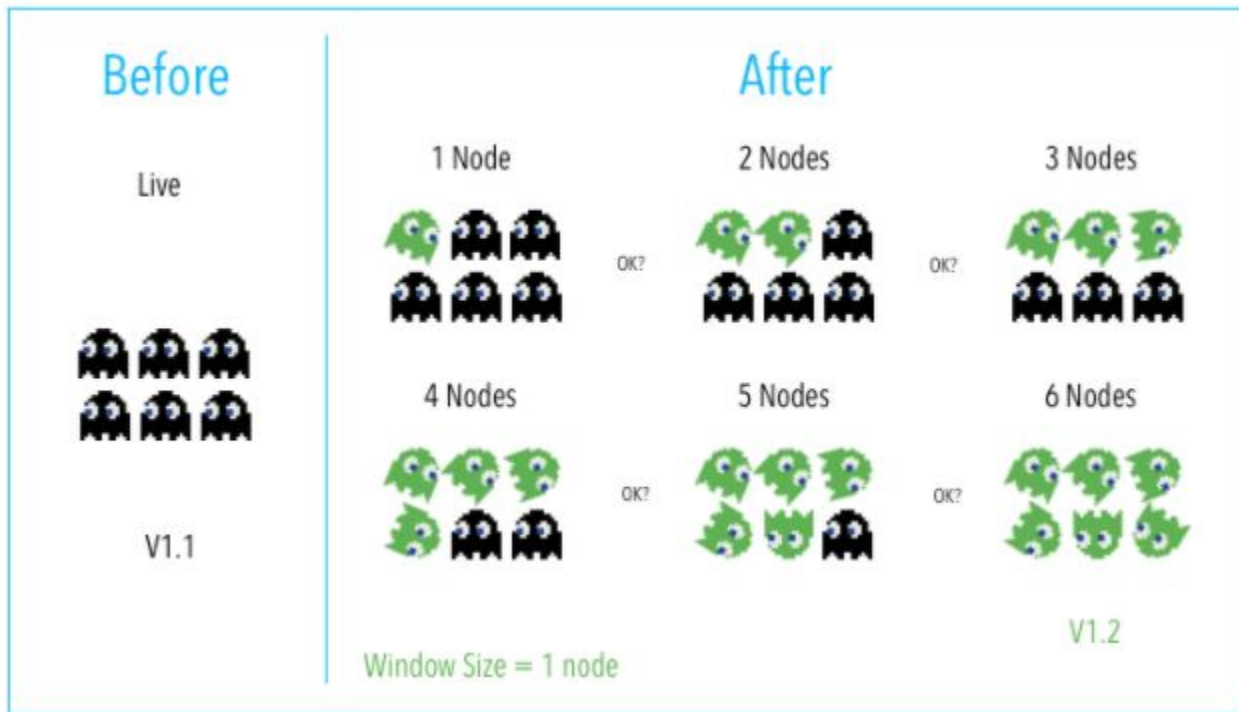
- ❶ The traffic block to specify the traffic distribution
- ❷ The unique name for this traffic block list item
- ❸ The Knative Revision that will participate in the traffic distribution
- ❹ The amount of traffic that the revision will receive; it is a numerical value in percentage



Deployment Strategies

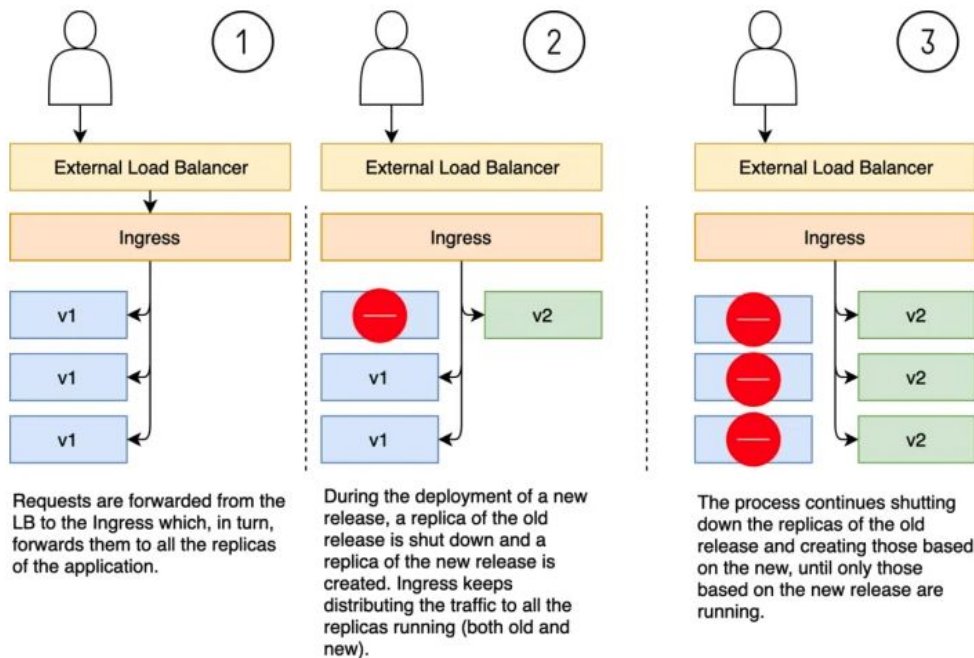


Rolling Deployment (Kubernetes default method)





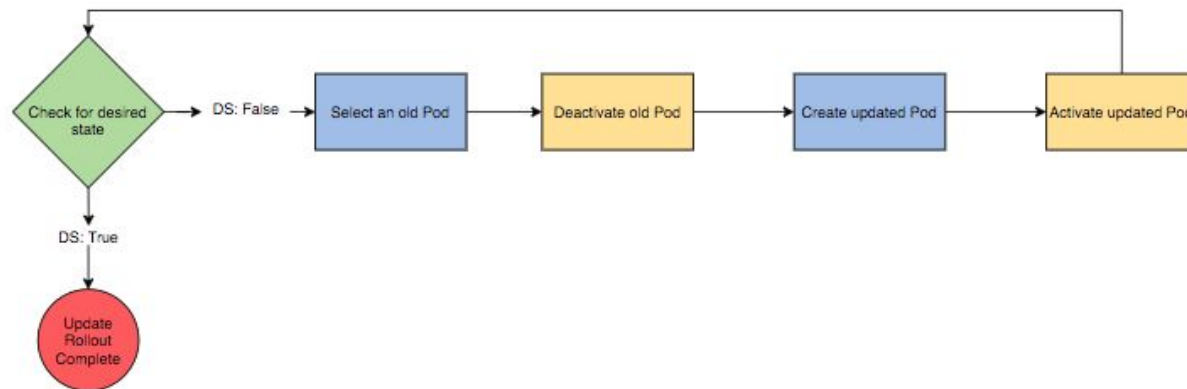
Rolling Updates Deployment Strategy





Rolling update strategy

The rolling update strategy is a gradual process that allows you to update your Kubernetes system with only a minor effect on performance and no downtime.



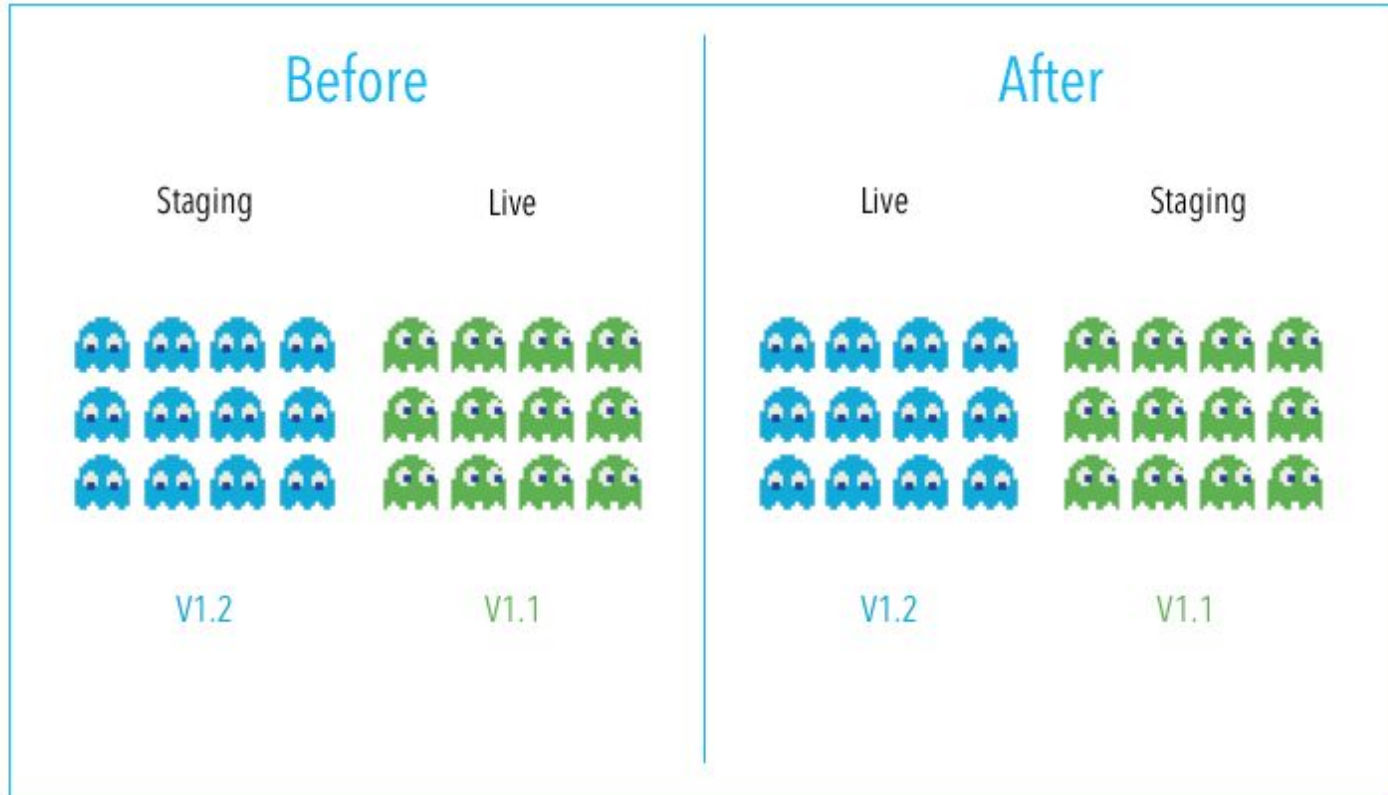


Blue Green Deployment

proprietary + confidential

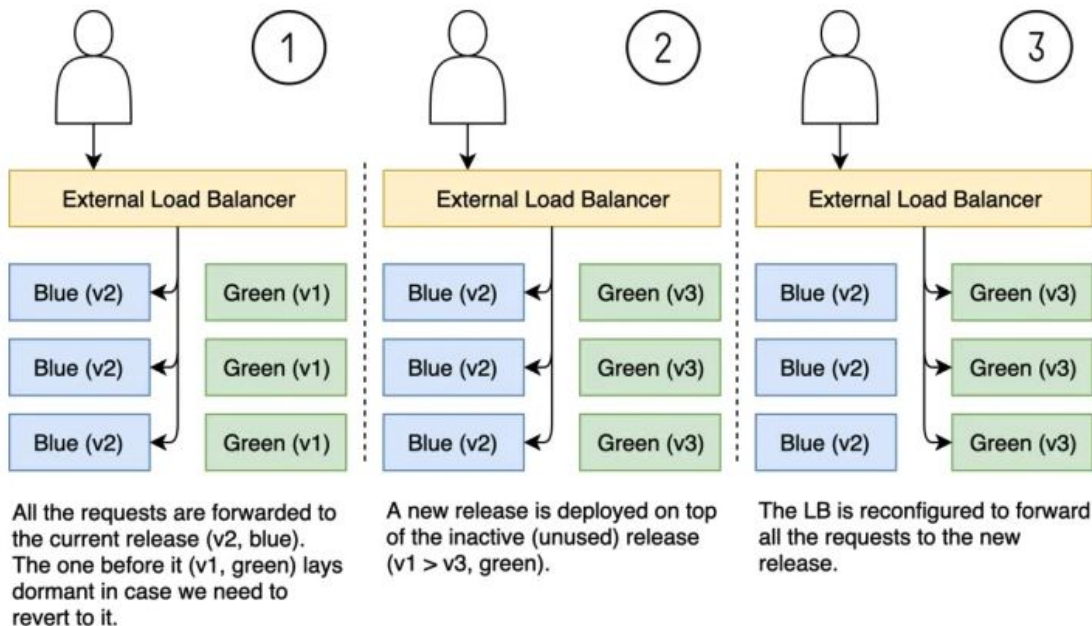
WIZELINE

www.wizeline.com





Blue-Green Deployments





Blue Green Knative Deployment

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: greeter
spec:
  template:
    metadata:
      name: greeter-v2
    spec:
      containers:
        - image: quay.io/rhdevelopers/knative-tutorial-greeter:quarkus
          env:
            - name: MESSAGE_PREFIX
              value: Namaste
          livenessProbe:
            httpGet:
              path: /healthz
          readinessProbe:
            httpGet:
              path: /healthz
      traffic:
        - tag: v1
          revisionName: greeter-v1
          percent: 100
        - tag: v2
          revisionName: greeter-v2
          percent: 0
        - tag: latest
          latestRevision: true
          percent: 0
```



Workshop Knative Traffic Management

Steps required.

```
kn service update blue-green-canary \  
  --image=quay.io/rhdevelopers/blue-green-canary \  
  --env BLUE_GREEN_CANARY_COLOR="#5bbf45" \  
  --env BLUE_GREEN_CANARY_MESSAGE="Namaste"
```

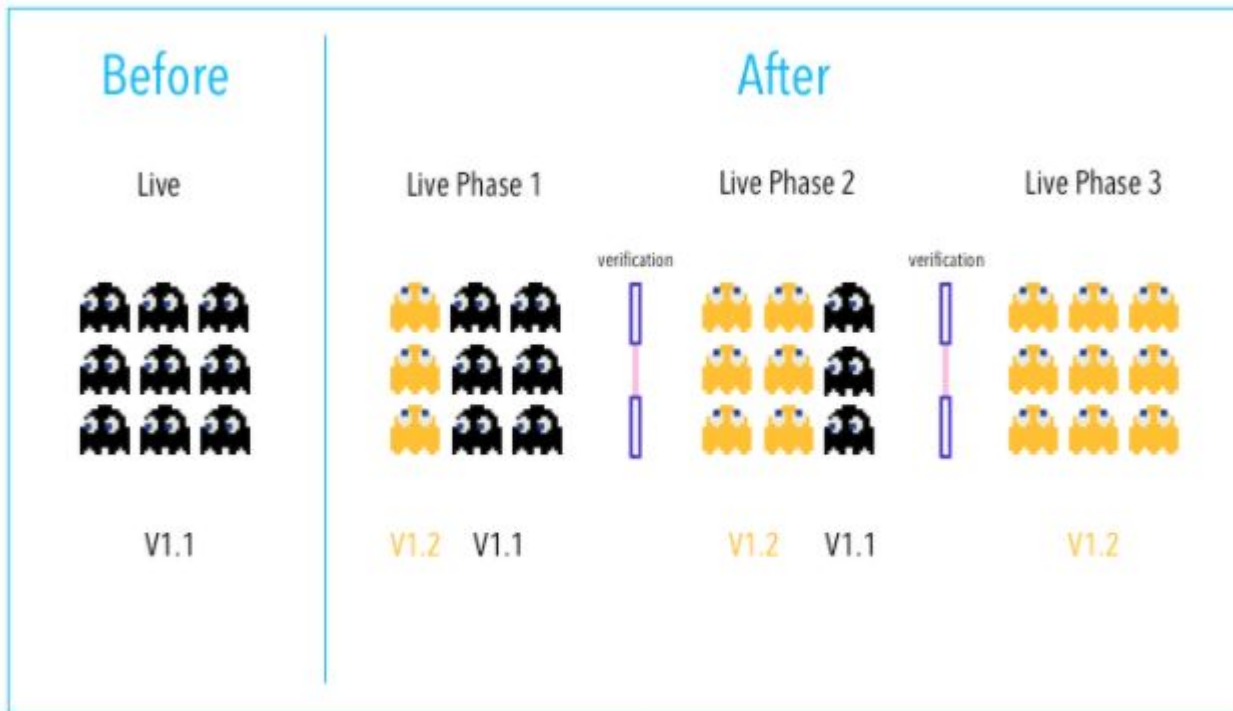


Canary Deployment

proprietary + confidential

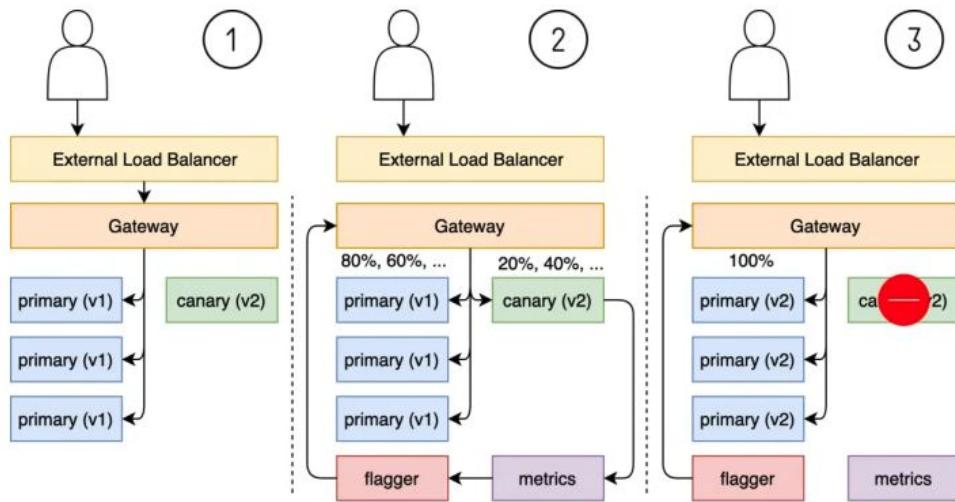
WIZELINE

www.wizeline.com





Canary deployment rollout



Requests are forwarded from the LB to the gateway which, in turn, forwards them to one of the Pods (e.g., round-robin).

When a new release (canary) is deployed, flagger configures the gateway to send portion of the requests to the new (canary) release and continue sending the rest to the old (primary). Metrics are collected in a database and Flagger uses them to decide whether to proceed with the deployment by increasing the percentage sent to canary and decreasing those that are sent to primary.

When the progressive rollout iterations are finished, primary is updated to the new release and canary is shut down. Flagger reconfigures the gateway so that all the traffic is forwarded to primary replicas.



Canary Knative Deployment

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: greeter
spec:
  template:
    metadata:
      name: greeter-v2
    spec:
      containers:
        - image: quay.io/rhdevelopers/knative-tutorial-greeter:quarkus
          env:
            - name: MESSAGE_PREFIX
              value: Namaste
          livenessProbe:
            httpGet:
              path: /healthz
          readinessProbe:
            httpGet:
              path: /healthz
      traffic:
        - tag: v1
          revisionName: greeter-v1
          percent: 80
        - tag: v2
          revisionName: greeter-v2
          percent: 20
        - tag: latest
          latestRevision: true
          percent: 0
```



Workshop Knative Traffic Management

proprietary + confidential

WIZELINE

www.wizeline.com

And for more examples →

<https://redhat-developer-demos.github.io/knative-tutorial/knative-tutorial/serving/traffic-distribution.html>